

AUTOMOBILE SUPPLY CHAIN MANAGEMENT SYSTEM

1st Devendra Mandava, dmandava
Computer Science and Engineering
University at Buffalo
Buffalo, United States
dmandava@buffalo.edu

2nd Hemanth Kongara, hkongara
Computer Science and Engineering
University at Buffalo
Buffalo, United States
hkongara@buffalo.edu

3rd Sethu Thakkilapati, sethutha
Computer Science and Engineering
University at Buffalo
Buffalo, United States
sethutha@buffalo.edu

Abstract—The Car Supply Chain Management System (CSCMS) serves as a specialized database solution designed to overcome the limitations of Excel in managing the complex automotive supply chain. This database system enhances operational efficiency by providing scalable data management, real-time updates, and collaborative features for multiple users. With relationships mapped between suppliers, products, and customers, the CSCMS ensures seamless transactional flow and data integrity. The adoption of the CSCMS is expected to streamline processes, reduce costs, and elevate customer satisfaction, thereby reinforcing the industry's competitiveness and adaptability.

Index Terms—Automotive supply chain, database system, operational efficiency, data management, real-time updates, collaborative features, data integrity, transactional flow, process streamlining, cost reduction, customer satisfaction, industry competitiveness, adaptability

I. INTRODUCTION

The Car Supply Chain Management System project aims to revolutionize the automotive industry by optimizing the end-to-end processes involved in sourcing, manufacturing, and delivering vehicles to customers. Leveraging a comprehensive database schema, this project meticulously organizes data into distinct tables for Suppliers, Products, Customers, Orders, Order Details, and Payment Information. Through efficient data management and insightful analytics, the project endeavors to enhance supply chain transparency, streamline operations, and ultimately elevate customer satisfaction.

II. PROBLEM STATEMENT

The automotive industry faces formidable challenges in managing its complex supply chain, characterized by multiple suppliers, diverse products, and intricate customer demands. Traditional methods of data management, such as Excel files, are inadequate for handling the voluminous and interconnected nature of supply chain data. This project aims to address these challenges by leveraging a comprehensive database solution to streamline supply chain operations, improve decision-making processes, and enhance overall efficiency.

III. BACKGROUND

The automotive supply chain is a multifaceted network involving numerous stakeholders, including suppliers, manufacturers, distributors, and customers. Managing this intricate ecosystem poses significant logistical and operational hurdles. Excel files, commonly used for data management in this industry, often lack the scalability, accessibility, and analytical capabilities required to effectively address these challenges. As a result, organizations struggle to gain meaningful insights from their data, leading to inefficiencies, delays, and missed opportunities.

IV. OBJECTIVE

The primary objective of this project is to develop a robust database infrastructure tailored to the unique needs of the automotive supply chain. By structuring data into distinct entities such as Suppliers, Products, Customers, Orders, and Payment Information, the project aims to provide a centralized repository for all supply chain-related data. This database will enable advanced analytics, predictive modeling, and real-time monitoring, empowering stakeholders to make data-driven decisions with confidence.

V. CONTRIBUTION TO THE PROBLEM DOMAIN

This project has the potential to revolutionize automotive supply chain management by offering a comprehensive solution to longstanding challenges. By transitioning from Excel files to a database-driven approach, organizations can unlock new possibilities for efficiency, agility, and innovation. The project's contribution lies in its ability to optimize processes, reduce costs, mitigate risks, and ultimately, enhance customer satisfaction. In an industry where competitiveness hinges on operational excellence, the impact of this project cannot be overstated.

VI. TARGET USERS

Automotive Manufacturers: Automotive manufacturers will use the database to manage their supply chain operations, track inventory levels, monitor production processes, and analyze

customer demand patterns. They will leverage the database to optimize production schedules, ensure timely delivery of vehicles, and maintain quality standards across their product lines.

Suppliers: Suppliers of automotive components, parts, and raw materials will utilize the database to coordinate orders, manage inventory levels, and track shipments. They will benefit from real-time visibility into demand forecasts, production schedules, and delivery timelines, enabling them to align their operations with customer requirements more effectively.

Distributors and Dealers: Distributors and dealerships will access the database to place orders, manage inventory at their respective locations, and track sales performance. They will use the database to forecast demand, optimize stocking levels, and analyze customer preferences to tailor their offerings accordingly.

Logistics and Transportation Providers: Logistics and transportation companies will leverage the database to optimize route planning, track shipments in transit, and ensure timely delivery of vehicles and components. They will utilize the database to streamline logistical operations, minimize transportation costs, and enhance overall supply chain efficiency.

VII. DATABASE ADMINISTRATION

The database will be administered by a dedicated team of IT professionals with expertise in database management and system administration. In a real-life scenario, this team may consist of database administrators (DBAs), system architects, network engineers, and cybersecurity specialists. Their responsibilities will include:

Database Design and Implementation: Designing the database schema, creating tables, defining relationships, and ensuring data integrity and normalization.

Data Management: Importing and exporting data, performing backups and restores, monitoring database performance, and optimizing query execution.

Security Management: Implementing access controls, encryption mechanisms, and authentication protocols to safeguard sensitive data from unauthorized access and cyber threats.

System Maintenance: Installing software updates, applying patches, tuning database parameters, and resolving technical issues to ensure the smooth operation of the database environment.

User Support and Training: Providing technical support to users, troubleshooting database-related issues, and conducting training sessions to educate users on database functionalities and best practices.

By having a dedicated team of database administrators, the organization can ensure the reliability, security, and scalability

of the database infrastructure, thereby maximizing its utility and value to the various stakeholders involved in the automotive supply chain.

Added new tables, updated ER diagram and database

VIII. ENTITY-RELATIONSHIP (E/R) DIAGRAM

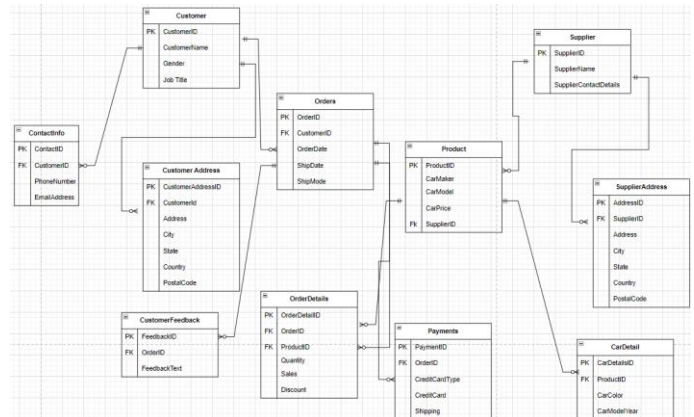


Fig. 1. ER Diagram

IX. DATA DESCRIPTION

Supplier's Data: Stores unique identifiers, names, and contact details for each supplier.

Supplier Address's Data: Contains addresses for suppliers, including detailed location information like city and postal code.

Product's Data (Car): Catalogs cars with details such as product ID, maker, model, price, and the associated supplier.

Car Detail's Data: Holds detailed specifications for each car, including color and model year.

Customer's Data: Records customer information including ID, name, gender, and job title.

Contact Info's Data: Captures contact details for each customer, such as phone numbers and email addresses.

Customer Address's Data: Provides address information for customers, including state and country.

Order's Data: Tracks order transactions, including dates, shipping mode, and linkage to customers.

Order Detail's Data: Details items within orders, covering aspects like quantity, sales, and discounts.

Payment's Data: Manages payment information related to orders, including credit card types and numbers.

Customer Feedback's Data: Collects customer feedback on orders, incorporating feedback text and related order ID.

IX. RELATIONSHIP BETWEEN TABLES

1.Supplier and SupplierAddress: One-to-Many, Each supplier can have multiple addresses, but each specific address (AddressID) is linked to only one supplier.

2.Supplier and Product (Car): One-to-Many, A supplier can supply multiple products (cars), but each car is provided by only one supplier.

3.Product (Car) and CarDetail: One-to-One, Each product (car) has one set of details specified in CarDetail, uniquely detailing aspects like color and model year for each car.

4.Customer and ContactInfo: One-to-Many, Each customer can have multiple contact information records, but each specific contact record is associated with only one customer.

5.Customer and CustomerAddress: One-to-Many, Customers can have multiple addresses registered, with each address uniquely linked to one customer.

6.Customer and Order: One-to-Many, Each customer can place multiple orders, but each order is associated with one specific customer.

7.Order and OrderDetail: One-to-Many, Each order can contain multiple items (OrderDetail), with each item's details, such as product ID and quantity, specifically linked to that order.

8.Order and Payment: One-to-One, Each order is associated with a unique payment record, detailing the transaction's payment method and costs.

9.Order and CustomerFeedback: One-to-One, Each order can receive one piece of feedback, directly linking customer feedback to the specific order it pertains to.

10.Product (Car) and OrderDetail: One-to-Many, Each car (product) can be part of multiple order details; that is, a car can be ordered multiple times in different orders or even within the same order as different order items.

X. ATTRIBUTES

Supplier's Data

- SupplierID INT (Primary Key): Uniquely identifies each supplier, non-null.
- SupplierName VARCHAR: Represents the name of the supplier, non-null.
- SupplierContactDetails TEXT: Holds contact details such as phone number and email, non-null.

SupplierAddress's Data

- AddressID INT (Primary Key): Uniquely identifies each address, non-null.
- SupplierID INT (Foreign Key to Supplier table): Links to the Supplier ID, non-null.
- Address TEXT: Street address of the supplier, non-null.
- City VARCHAR: City of the supplier's address, non-null.
- State VARCHAR: State of the supplier's address, non-null.
- Country VARCHAR: Country of the supplier's address, non-null.
- PostalCode INT: Postal code of the supplier's address, non-null.

Product's Data (Car)

- ProductID INT (Primary Key): Uniquely defines the product, non-null.

- CarMaker VARCHAR: Represents the maker of the car, non-null.
- CarModel VARCHAR: Represents the model of the car, non-null.
- CarPrice FLOAT: Represents the price of the car, non-null.
- SupplierID INT (Foreign Key to Supplier table): Links to the supplier of the car, non-null.

CarDetail's Data

- CarDetailID INT (Primary Key): Uniquely identifies car detail records, non-null.
- ProductID INT (Foreign Key to Product table): Links to the car product, non-null.
- CarColor VARCHAR: Represents the color of the car, non-null.
- CarModelYear INT: Represents the model year of the car, non-null.

Customer's Data

- CustomerID INT (Primary Key): Uniquely identifies each customer, non-null.
- CustomerName VARCHAR: Represents the full name of the customer, non-null.
- Gender CHAR(1): Represents the gender of the customer, non-null.
- JobTitle VARCHAR: Represents the job title of the customer, non-null.

ContactInfo's Data

- ContactID INT (Primary Key): Uniquely identifies contact information, non-null.
- CustomerID INT (Foreign Key to Customer table): Links to the customer, non-null.
- PhoneNumber VARCHAR: Represents the phone number of the customer, non-null.
- EmailAddress VARCHAR: Represents the email address of the customer, non-null.

CustomerAddress's Data

- CustomerAddressID INT (Primary Key): Uniquely identifies customer addresses, non-null.
- CustomerID INT (Foreign Key to Customer table): Links to the customer, non-null.
- Address TEXT: Street address of the customer, non-null.
- City VARCHAR: City of the customer's address, non-null.
- State VARCHAR: State of the customer's address, non-null.
- Country VARCHAR: Country of the customer's address, non-null.
- PostalCode INT: Postal code of the customer's address, non-null.

Order's Data

- OrderID INT (Primary Key): Uniquely defines each order, non-null.
- CustomerID INT (Foreign Key to Customer table): Links to the customer who made the order, non-null.
- OrderDate DATE: Represents the date the order was placed, non-null.
- ShipDate DATE: Represents the date the order is shipped,

non-null.

- ShipMode VARCHAR: Represents the mode of shipping used for the order, non-null.

OrderDetail's Data

- OrderDetailID INT (Primary Key): Uniquely defines each order item, non-null.
- OrderID INT (Foreign Key to Order table): Links to the order, non-null.
- ProductID INT (Foreign Key to Product table): Links to the product ordered, non-null.
- Quantity INT: Represents the number of products ordered, non-null.
- Sales FLOAT: Represents the total sales amount, non-null.
- Discount FLOAT: Represents the discount given on the order item, non-null.

Payment's Data

- PaymentID INT (Primary Key): Uniquely defines each payment record, non-null.
- OrderID INT (Foreign Key to Order table): Links to the order, non-null.
- CreditCardType VARCHAR: Represents the type of credit card used, non-null.
- CreditCard VARCHAR: Represents the credit card number, non-null.
- Shipping FLOAT: Represents the shipping cost, non-null.

CustomerFeedback's Data

- FeedbackID INT (Primary Key): Uniquely identifies each feedback entry, non-null.
- OrderID INT (Foreign Key to Order table): Links to the order, non-null.
- FeedbackText TEXT: Represents the text of the feedback provided by the customer, non-null.

X. ACTIONS ON FOREIGN KEY WHEN PRIMARY KEY DELETED:

No Action (or Restrict): This is the default behavior in many databases. If there are any foreign key dependencies, the system will prevent the deletion of the primary key. This ensures that no orphan records are left in the child tables.

Cascade: When the primary key is deleted, a cascade action will automatically delete all related records in the foreign key table that depend on the primary key. This action ensures that there are no dangling references in the database and helps maintain referential integrity by automatically removing all dependent entities.

Set Null: This action sets the foreign key value in the child table to NULL when the referenced primary key is deleted. It's useful when it's acceptable to retain the dependent records without linking them to the parent record, effectively nullifying the relationship but retaining the child records.

Set Default: Similar to the Set Null action, but instead of setting the foreign key to NULL, it sets it to a default value specified in the table schema. This can be used to reassign dependent records to a default value when the primary entity is removed.

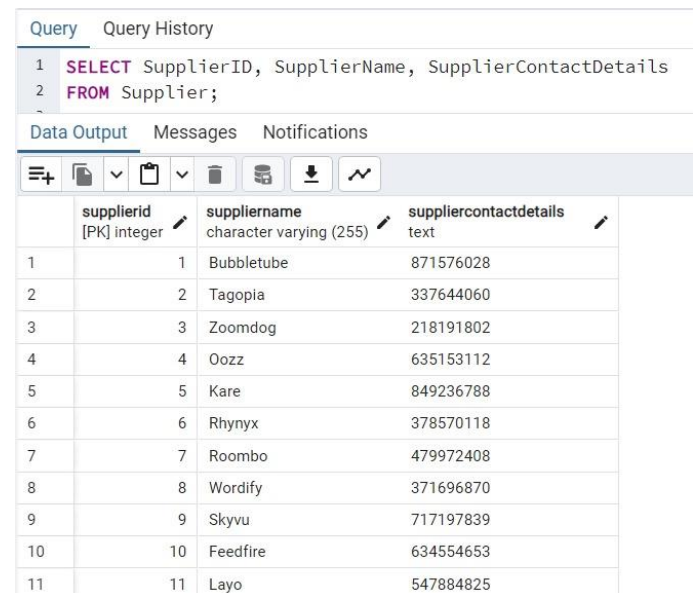
No Action vs. Restrict: While these are often considered the same, depending on the database system, "Restrict" might be an explicitly declared action that prevents the deletion of a primary key if there are any dependencies, whereas "No Action" could potentially defer the integrity check to the end of the transaction.

XI. APPROACH USED TO CREATE TABLES FROM CSV

In our approach to creating database tables from CSV files, we initially preprocessed the data using Python to segment and structure it into distinct tables. This preprocessing involved cleaning, validating, and organizing the data according to the schema requirements of our target database. Following the preprocessing phase, we manually loaded the structured data into PostgreSQL, utilizing its robust SQL capabilities to ensure each dataset was accurately inserted into its respective table. This method allowed for precise control over the data import process, ensuring data integrity and consistency across our database system.

XII. PERFORMED QUERIES:

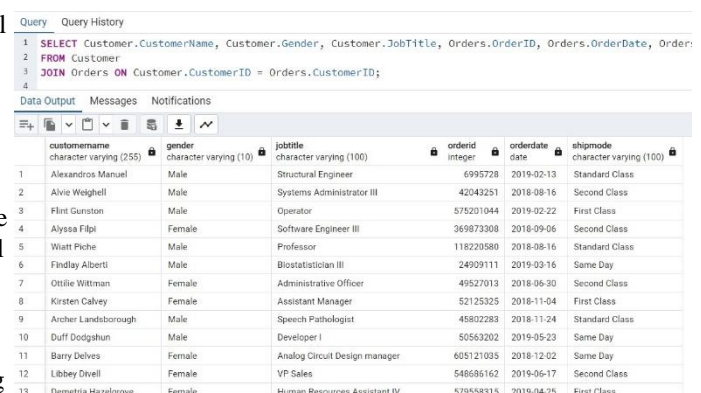
Query 1 : Retrieve All Suppliers



	supplierid [PK] integer	suppliername character varying (255)	suppliercontactdetails text
1	1	Bubbletube	871576028
2	2	Tagopia	337644060
3	3	Zoomdog	218191802
4	4	Oozz	635153112
5	5	Kare	849236788
6	6	Rhynyx	378570118
7	7	Roombo	479972408
8	8	Wordify	371696870
9	9	Skyvu	717197839
10	10	Feedfire	634554653
11	11	Layo	547884825

Fig.2.

Query 2 : Customer Information and Their Orders



	customername character varying (255)	gender character varying (10)	jobtitle character varying (100)	orderid integer	orderdate date	shipmode character varying (100)
1	Alexandros Manuel	Male	Structural Engineer	6995728	2019-02-13	Standard Class
2	Alvie Weighell	Male	Systems Administrator III	42043251	2018-08-16	Second Class
3	Flint Gunston	Male	Operator	575201944	2019-02-22	First Class
4	Alyssa Filpi	Female	Software Engineer III	369873308	2018-09-06	Second Class
5	Watt Piche	Male	Professor	118220580	2018-08-16	Standard Class
6	Findlay Alberti	Male	Biostatistician III	24909111	2019-02-16	Same Day
7	Ottile Wittman	Female	Administrative Officer	49527013	2018-06-30	Second Class
8	Kirsten Calvey	Female	Assistant Manager	52125325	2018-11-04	First Class
9	Archer Landsborough	Male	Speech Pathologist	45802283	2018-11-24	Standard Class
10	Duff Doddgshun	Male	Developer I	50563202	2019-05-23	Same Day
11	Barry Delves	Female	Analog Circuit Design manager	605121035	2018-12-02	Same Day
12	Libbey Driell	Female	VP Sales	548686162	2019-06-17	Second Class
13	Demetria Hazelgrove	Female	Human Resources Assistant IV	579558315	2019-04-25	First Class

Fig.3.

Query 3: Detailed Car Information

Query

Query History

1

SELECT Product.CarMaker, Product.CarModel, Product.CarPrice, CarDetail.CarColor, CarDetail.CarModelYear

2

FROM Product

3

JOIN CarDetail ON Product.ProductID = CarDetail.ProductID;

4

Data Output

Messages

Notifications

	car maker character varying (100)	car model character varying (100)	car price numeric	car color character varying (50)	car model year integer
1	Dodge	Ram 2500	521963.45	Goldenrod	2007
2	Toyota	Tundra	672222.04	Crimson	2010
3	GMC	Savana 1500	504465.72	Crimson	2011
4	Volkswagen	Cabriolet	646077.11	Fuscia	1990
5	Mercury	Mariner	699890.24	Teal	2009
6	Toyota	Land Cruiser	694873.94	Crimson	2005
7	Subaru	Impreza	546977.55	Indigo	1999
8	BMW	X6	529353.8	Purple	2013
9	Mitsubishi	Galant	706964.35	Orange	2011
10	Subaru	Justy	798573.59	Crimson	1988
11	Mercedes-Benz	M-Class	722114.57	Mauv	2002
12	Ford	Taurus X	761182.85	Crimson	2008
13	Mercedes-Benz	CL-Class	568015.89	Red	2011

Fig.4.

Query 4: Orders and Associated Payment Information

Query

Query History

1

SELECT Orders.OrderID, Orders.OrderDate, Payment.CreditCardType, Payment.CreditCard, Payment.Shipping

2

FROM Orders

3

JOIN Payment ON Orders.OrderID = Payment.OrderID;

4

Data Output

Messages

Notifications

orderid

integer

creditcardtype

character varying (50)

creditcard

character varying (20)

shipping

character varying (20)

1

6995728

2019-02-13

jcb

3549221112237767

Truck

2

42043251

2018-08-16

jcb

3557159608180902

Air

3

575201044

2019-02-22

jcb

3529909223663921

Truck

4

369873308

2018-09-06

china-unionpay

5602235978541517

Air

5

118220580

2018-08-16

laser

6706247841149207

Air

6

24909111

2019-03-16

diners-club-carte-blanche

30528106894581

Truck

7

49527013

2018-06-30

jcb

3589144247894735

Truck

8

52125325

2018-11-04

jcb

3557048508901444

Air

9

45802283

2018-11-24

jcb

3534727482148019

Truck

10

50563202

2019-05-23

diners-club-enroute

201441036983941

Air

11

605121035

2018-12-02

jcb

3558313957113134

Truck

12

548886162

2019-06-17

diners-club-enroute

201504347329142

Truck

13

579558315

2019-04-25

jcb

3542457262271746

Air

Fig.5.

Query 5: Retrieve Customer Feedback with Order Details

Query

Query History

1

SELECT

2

CustomerFeedback.FeedbackID,

3

CustomerFeedback.FeedbackText,

4

Orders.OrderID,

5

Orders.OrderDate,

6

Orders.CustomerID

7

FROM

8

CustomerFeedback

9

JOIN

10

Orders ON CustomerFeedback.OrderID = Orders.OrderID;

Data Output

Messages

Notifications

feedbackid

integer

feedbacktext

text

orderid

integer

orderdate

date

customerid

integer

1

464898

Good

6995728

2019-02-13

67457594

2

874500

Okay

42043251

2018-08-16

58411135

3

548869

Very Bad

575201044

2019-02-22

5915307

4

581640

Bad

369873308

2018-09-06

51655189

5

745482

Good

118220580

2018-08-16

658110001

6

768011

Very Bad

24909111

2019-03-16

31722328

7

874507

Bad

49527013

2018-06-30

54162018

8

735245

Okay

52125325

2018-11-04

449110060

9

469007

Very Bad

45802283

2018-11-24

548685732

10

182293

Very Bad

50563202

2019-05-23

9551015

11

782360

Okay

605121035

2018-12-02

54838101

12

909336

Good

548686162

2019-06-17

49825119

Fig.6.

XIII. BCNF PROOF:

To ensure that the schema for the Car Supply Chain Management system is in Boyce-Codd Normal Form (BCNF), we must confirm that every non-trivial functional dependency has its left-hand side as a superkey. Here's an analysis of each table based on the functional dependencies.

Supplier

- **Functional Dependency:** SupplierID → SupplierName, SupplierContactDetails

- **Analysis:** SupplierID is the primary key and uniquely determines other attributes in the table, qualifying as a superkey. This table satisfies the BCNF condition.

SupplierAddress

- **Functional Dependency:** AddressID → SupplierID, Address, City, State, Country, PostalCode

- **Analysis:** AddressID is the primary key, making it a superkey that uniquely determines all other attributes in the table. This table is in BCNF.

Product (Car)

- **Functional Dependency:** ProductID → CarMaker, CarModel, CarPrice, SupplierID

- **Analysis:** ProductID, as the primary key, functions as a superkey and uniquely determines the specified attributes, adhering to BCNF requirements.

CarDetail

- **Functional Dependency:** CarDetailID → ProductID, CarColor, CarModelYear

- **Analysis:** CarDetailID is the primary key and a superkey that uniquely determines other related attributes. The table is in BCNF.

Customer

- **Functional Dependency:** CustomerID → CustomerName, Gender, JobTitle

- **Analysis:** CustomerID is the primary key and satisfies the BCNF condition by uniquely determining all other non-key attributes in the table.

ContactInfo

- **Functional Dependency:** ContactID → CustomerID, PhoneNumber, EmailAddress

- **Analysis:** ContactID, as the primary key, qualifies as a superkey, ensuring the table meets BCNF standards.

CustomerAddress

- **Functional Dependency:** CustomerAddressID → CustomerID, Address, City, State, Country, PostalCode

- **Analysis:** CustomerAddressID is the primary key and a superkey, ensuring the table adheres to the BCNF form.

Order

- **Functional Dependency:** OrderID → CustomerID, OrderDate, ShipDate, ShipMode

- **Analysis:** OrderID is the primary key and acts as a superkey, satisfying the BCNF criteria.

OrderDetail

- **Functional Dependency:** OrderDetailID → OrderID, ProductID, Quantity, Sales, Discount

- **Analysis:** OrderDetailID is the primary key and functions as a superkey, making the table BCNF compliant.

Payment

- **Functional Dependency:** PaymentID → OrderID, CreditCardType, CreditCard, Shipping

- **Analysis:** PaymentID is the primary key and superkey, ensuring this table is in BCNF.

CustomerFeedback

- **Functional Dependency:** FeedbackID → OrderID, FeedbackText

- **Analysis:** FeedbackID is the primary key, serving as a superkey. This table meets the BCNF conditions.

Each table's functional dependencies confirm that the primary key functions as a superkey, which functionally determines all other attributes in the table. These conditions satisfy the requirements for BCNF in all tables, eliminating the need for further decomposition to address anomalies or integrity issues. The Entity-Relationship (E/R) diagram thus remains unchanged from its initial state, reflecting a well-structured relational database design that adheres to BCNF criteria.

XIV. USE OF INDEXING:

When dealing with a larger dataset, a major challenge faced was the increased retrieval times for queries, particularly those that required grouping operations. As the size of the dataset grew, the duration required to sift through the data and extract pertinent information also rose markedly. To tackle this problem, indexing techniques were implemented to enhance the efficiency of data retrieval. The time it took to execute a query before the introduction of indexing is illustrated in figure 7

Query:1

Query	Query History
1	<code>explain analyze</code>
2	<code>SELECT p.CarMaker, p.CarModel, SUM(od.Quantity) AS TotalQuantity</code>
3	<code>FROM OrderDetail od</code>
4	<code>JOIN Orders o ON od.OrderID = o.OrderID</code>
5	<code>JOIN Product p ON od.ProductID = p.ProductID</code>
6	<code>WHERE o.OrderDate >= (</code>
7	<code>SELECT DATE_TRUNC('month', MAX(OrderDate))</code>
8	<code>FROM Orders</code>
9	<code>) AND o.OrderDate < (</code>
10	<code>SELECT (DATE_TRUNC('month', MAX(OrderDate)) + INTERVAL '1 month')</code>
11	<code>FROM Orders</code>
12	<code>)</code>
13	<code>GROUP BY p.CarMaker, p.CarModel</code>
14	<code>ORDER BY TotalQuantity DESC;</code>
Data Output Messages Explain X Notifications	
QUERY PLAN	
text	
19	-> Hash (cost=20.08..20.08 rows=4 width=4) (actual time=0.461..0.462 rows=46 loops=1)
20	Buckets: 1024 Batches: 1 Memory Usage: 10kB
21	-> Seq Scan on orders o (cost=0.00..20.08 rows=4 width=4) (actual time=0.292..0.438 rows=46 loops=1)
22	Filter: ((orderdate >= \$0) AND (orderdate < \$1))
23	Rows Removed by Filter: 826
24	-> Index Scan using product_pkey on product p (cost=0.28..0.35 rows=1 width=18) (actual time=0.004..0.004 rows=1 loops=1)
25	Index Cond: (productid = od.productid)
26	Planning Time: 0.352 ms
27	Execution Time: 1.065 ms

Fig.7.

To speed up query performance on the 'orderdate' column, the command `CREATE INDEX idx_orderdate ON orders(orderdate);` was used to create an index.

Query	Query History
1	<code>explain analyze</code>
2	<code>SELECT p.CarMaker, p.CarModel, SUM(od.Quantity) AS TotalQuantity</code>
3	<code>FROM OrderDetail od</code>
4	<code>JOIN Orders o ON od.OrderID = o.OrderID</code>
5	<code>JOIN Product p ON od.ProductID = p.ProductID</code>
6	<code>WHERE o.OrderDate >= (</code>
7	<code>SELECT DATE_TRUNC('month', MAX(OrderDate))</code>
8	<code>FROM Orders</code>
9	<code>) AND o.OrderDate < (</code>
10	<code>SELECT (DATE_TRUNC('month', MAX(OrderDate)) + INTERVAL '1 month')</code>
11	<code>FROM Orders</code>
12	<code>)</code>
13	<code>GROUP BY p.CarMaker, p.CarModel</code>
Data Output Messages Explain X Notifications	
QUERY PLAN	
text	
28	Buckets: 1024 Batches: 1 Memory Usage: 10kB
29	-> Bitmap Heap Scan on orders o (cost=4.32..11.30 rows=4 width=4) (actual time=0.098..0.109 rows=46 loops=1)
30	Recheck Cond: ((orderdate >= \$1) AND (orderdate < \$3))
31	Heap Blocks: exact=7
32	-> Bitmap Index Scan on idx_orderdate (cost=0.00..4.32 rows=4 width=0) (actual time=0.092..0.092 rows=46 loops=1)
33	Index Cond: ((orderdate >= \$1) AND (orderdate < \$3))
34	-> Index Scan using product_pkey on product p (cost=0.28..0.35 rows=1 width=18) (actual time=0.001..0.001 rows=1 loops=46)
35	Index Cond: (productid = od.productid)
36	Planning Time: 2.382 ms
37	Execution Time: 0.513 ms

Fig.8.

Query:2

Query	Query History
1	<code>explain analyze</code>
2	<code>SELECT s.SupplierName, p.CarModel, p.CarPrice</code>
3	<code>FROM Product p</code>
4	<code>JOIN Supplier s ON p.SupplierID = s.SupplierID</code>
5	<code>WHERE p.CarPrice > (SELECT AVG(CarPrice) FROM Product)</code>
6	<code>ORDER BY p.CarPrice DESC;</code>
Data Output Messages Explain X Notifications	
QUERY PLAN	
text	
1	Sort (cost=72.26..72.99 rows=291 width=22) (actual time=1.039..1.063 rows=442 loops=1)
2	Sort Key: p.carprice DESC
3	Sort Method: quicksort Memory: 59kB
4	InitPlan 1 (returns \$0)
5	-> Aggregate (cost=17.90..17.91 rows=1 width=32) (actual time=0.249..0.250 rows=1 loops=1)
6	-> Seq Scan on product (cost=0.00..15.72 rows=872 width=8) (actual time=0.007..0.084 rows=872 loops=1)
7	-> Hash Join (cost=21.54..42.44 rows=291 width=22) (actual time=0.618..0.872 rows=442 loops=1)
8	Hash Cond: (s.supplierid = p.supplierid)
9	-> Seq Scan on supplier s (cost=0.00..14.72 rows=872 width=11) (actual time=0.012..0.104 rows=872 loops=1)
10	-> Hash (cost=17.90..17.90 rows=291 width=19) (actual time=0.592..0.592 rows=442 loops=1)
11	Buckets: 1024 Batches: 1 Memory Usage: 32kB
12	-> Seq Scan on product p (cost=0.00..17.90 rows=291 width=19) (actual time=0.262..0.500 rows=442 loops=1)
13	Filter: (carprice > \$0)
14	Rows Removed by Filter: 430
15	Planning Time: 0.244 ms
16	Execution Time: 1.122 ms

Fig.9.

To optimize query performance, three indexes were created: `CREATE INDEX idx_product_carprice ON product(carprice);` to facilitate faster searches based on car prices, `CREATE INDEX idx_supplier_supplierid ON supplier(supplierid);` to enhance lookups by supplier ID, and `CREATE INDEX idx_product_supplierid ON product(supplierid);` to improve access times when querying products by their supplier ID.

Query

Query History

```

1  explain analyze
2  SELECT s.SupplierName, p.CarModel, p.CarPrice
3  FROM Product p
4  JOIN Supplier s ON p.SupplierID = s.SupplierID
5  WHERE p.CarPrice > (SELECT AVG(CarPrice) FROM Product)
6  ORDER BY p.CarPrice DESC;

```

Data Output
Messages
Explain X
Notifications

+

⌵

⌵

⌵

⌵

⌵

⌵

⌵

+

⌵

⌵

⌵

⌵

⌵

⌵

⌵

QUERY PLAN

text

1 Sort (cost=72.26..72.99 rows=291 width=22) (actual time=0.628..0.642 rows=442 loops=1)
2 Sort Key: p.carprice DESC
3 Sort Method: quicksort Memory: 59kB
4 InitPlan 1 (returns \$0)
5 -> Aggregate (cost=17.90..17.91 rows=1 width=32) (actual time=0.154..0.154 rows=1 loops=1)
6 -> Seq Scan on product (cost=0.00..15.72 rows=872 width=8) (actual time=0.004..0.051 rows=872 loops=1)
7 -> Hash Join (cost=21.54..42.44 rows=291 width=22) (actual time=0.368..0.516 rows=442 loops=1)
8 Hash Cond: (s.supplierid = p.supplierid)
9 Seq Scan on supplier s (cost=0.00..14.72 rows=872 width=11) (actual time=0.009..0.053 rows=872 loops=1)
10 Hash (cost=17.90..17.90 rows=291 width=19) (actual time=0.354..0.354 rows=442 loops=1)
11 Buckets: 1024 Batches: 1 Memory Usage: 32kB
12 -> Seq Scan on product p (cost=0.00..17.90 rows=291 width=19) (actual time=0.162..0.301 rows=442 loops=1)
13 Filter: (carprice > \$0)
14 Rows Removed by Filter: 430
15 Planning Time: 2.777 ms
16 Execution Time: 0.671 ms

Query:3

```
Query    Query History
1  explain analyze
2  SELECT p.CarModel, SUM(od.Sales) AS TotalSales
3  FROM OrderDetail od
4  JOIN Product p ON od.ProductID = p.ProductID
5  GROUP BY p.CarModel
6  ORDER BY TotalSales DESC;
```

Data Output Messages Explain X Notifications

QUERY PLAN
text

```
1  Sort (cost=75.93..77.11 rows=473 width=39) (actual time=1.556..1.571 rows=473 loops=1)
   Sort Key: (sum(od.sales)) DESC
3  Sort Method: quicksort  Memory: 54kB
4  -> HashAggregate (cost=49.00..54.91 rows=473 width=39) (actual time=1.323..1.452 rows=473 loops=1)
   Group Key: p.carmodel
5  -> Hash Join (cost=26.62..44.64 rows=872 width=15) (actual time=0.298..0.756 rows=872 loops=1)
   Hash Cond: (od.productid = p.productid)
6  -> Seq Scan on orderdetail od (cost=0.00..15.72 rows=872 width=12) (actual time=0.029..0.157 rows=872 loops=1)
   -> Hash (cost=15.72..15.72 rows=872 width=11) (actual time=0.261..0.262 rows=872 loops=1)
   Buckets: 1024 Batches: 1 Memory Usage: 47kB
7  -> Seq Scan on product p (cost=0.00..15.72 rows=872 width=11) (actual time=0.011..0.126 rows=872 loops=1)
12 Planning Time: 0.223 ms
13 Execution Time: 1.642 ms
```

To improve query efficiency, the commands ``CREATE INDEX idx_product_productid ON product(productid);`` and ``CREATE INDEX idx_orderdetail_productid ON orderdetail(productid);`` were used to create indexes, enhancing data retrieval based on product IDs in both the 'product' and 'orderdetail' tables.

Query

Query History

```

1 explain analyze
2 SELECT p.CarModel, SUM(od.Sales) AS TotalSales
3 FROM OrderDetail od
4 JOIN Product p ON od.ProductID = p.ProductID
5 GROUP BY p.CarModel
6 ORDER BY TotalSales DESC;

```

Data Output

Messages

Explain X

Notifications

+

📄

📄

🗑️

🗑️

📄

📄

QUERY PLAN

text

1

Sort (cost=75.93..77.11 rows=473 width=39) (actual time=0.981..0.993 rows=473 loops=1)

2

Sort Key: (sum(od.sales)) DESC

3

Sort Method: quicksort Memory: 54kB

4

→ HashAggregate (cost=49.00..54.91 rows=473 width=39) (actual time=0.710..0.859 rows=473 loops=1)

5

Group Key: p.carmodel

6

→ Hash Join (cost=26.62..44.64 rows=872 width=15) (actual time=0.184..0.400 rows=872 loops=1)

7

Hash Cond: (od.productid = p.productid)

8

→ Seq Scan on orderdetail od (cost=0.00..15.72 rows=872 width=12) (actual time=0.010..0.085 rows=872 loops=1)

9

→ Hash (cost=15.72..15.72 rows=872 width=11) (actual time=0.167..0.167 rows=872 loops=1)

10

Buckets: 1024 Batches: 1 Memory Usage: 47kB

11

→ Seq Scan on product p (cost=0.00..15.72 rows=872 width=11) (actual time=0.008..0.074 rows=872 loops=1)

12

Planning Time: 0.250 ms

13

Execution Time: 1.062 ms

The execution time taken for a query after use of indexing can be addressed in Fig.8, Fig.10, Fig.12.

Overall, by addressing the retrieval time issue through indexing, able to improve the performance of the database when handling larger datasets, ensuring that queries were executed more efficiently.

XV. QUERYING:

Query1: Insert a new supplier and their address:

```
Query    Query History
1  INSERT INTO Supplier (SupplierName, SupplierContactDetails)
2  VALUES ('Quality Goods Inc', 'contact@qualitygoods.com');
3
4  INSERT INTO SupplierAddress (SupplierID, SupplierAddress, City, State, Country, PostalCode)
5  VALUES ((SELECT MAX(supplierid) FROM Supplier),
6           '4567 Quality Ave', 'Quality Town', 'Quality State', 'Quality Country', '67890');
7
```

Data Output Messages Notifications

INSERT 0 1

Query returned successfully in 39 msec.

Query 2: Insert a new product provided by the last added supplier

Query	Query History
1	INSERT INTO Product (CarMaker, CarModel, CarPrice, SupplierID)
2	VALUES ('Toyota', 'Corolla', 20000.00, (SELECT MAX(SupplierID) FROM Supplier));
3	

Data Output	Messages	Notifications
INSERT 0 1		
Query returned successfully in 45 msec.		

Query3: Deleting a Customer and All Related Records

Query	Query History
1	BEGIN;
2	DELETE FROM CustomerFeedback
3	WHERE OrderID IN (
4	SELECT OrderID FROM Orders WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche'));
5	DELETE FROM Payment
6	WHERE OrderID IN (
7	SELECT OrderID FROM Orders WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche'));
8	DELETE FROM OrderDetail
9	WHERE OrderID IN (
10	SELECT OrderID FROM Orders WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche'));
11	DELETE FROM Orders
12	WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche');
13	DELETE FROM CustomerAddress
14	WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche');
15	DELETE FROM ContactInfo
16	WHERE CustomerID = (SELECT CustomerID FROM Customer WHERE CustomerName = 'Louie Hinsche');
17	DELETE FROM Customer
18	WHERE CustomerName = 'Louie Hinsche';
19	
20	COMMIT;
21	

Data Output	Messages	Notifications
COMMIT		
Query returned successfully in 43 msec.		

Query 4 : Deleting a Supplier and All Related Records

Query	Query History
1	BEGIN;
2	DELETE FROM CarDetail
3	WHERE ProductID IN (
4	SELECT ProductID FROM Product WHERE SupplierID = (
5	SELECT SupplierID FROM Supplier WHERE SupplierName = 'Quality Goods Inc'))
6);
7	DELETE FROM Product
8	WHERE SupplierID = (SELECT SupplierID FROM Supplier WHERE SupplierName = 'Quality Goods Inc');
9	
10	DELETE FROM SupplierAddress
11	WHERE SupplierID = (SELECT SupplierID FROM Supplier WHERE SupplierName = 'Quality Goods Inc');
12	
13	DELETE FROM Supplier
14	WHERE SupplierName = 'Quality Goods Inc';
15	
16	COMMIT;

Data Output	Messages	Notifications
COMMIT		
Query returned successfully in 46 msec.		

Query 5: Update the price of a specific model from a specific maker

Query	Query History
1	UPDATE Product
2	SET CarPrice = CarPrice * 1.1 -- Increase price by 10%
3	WHERE CarMaker = 'Toyota' AND CarModel = 'Corolla';
4	

Data Output	Messages	Notifications
UPDATE 3		
Query returned successfully in 49 msec.		

Query 6: Update contact details for a specific customer

Query	Query History
1	UPDATE ContactInfo
2	SET EmailAddress = 'newemail22@domain.com', PhoneNumber = '1234597890'
3	WHERE CustomerID = (SELECT CustomerID
4	FROM Customer WHERE CustomerName = 'Flint Gunston');
5	

Data Output	Messages	Notifications
UPDATE 1		
Query returned successfully in 53 msec.		

Query7: Join multiple tables to get a detailed order list including customer info, order details, and product info

Query

Query History

1

SELECT c.CustomerName, o.OrderDate, p.CarMaker, p.CarModel, od.Quantity, od.Sales

2

FROM Orders o

3

JOIN Customer c ON o.CustomerID = c.CustomerID

4

JOIN OrderDetail od ON o.OrderID = od.OrderID

5

JOIN Product p ON od.ProductID = p.ProductID

6

WHERE o.OrderDate BETWEEN '2019-01-01' AND '2019-5-31'

7

ORDER BY o.OrderDate DESC;

8

Data Output

Messages

Notifications

<

Query8: Aggregate function to get total sales per car model grouped by model

Query	Query History
1	SELECT p.CarModel, SUM(od.Sales) AS TotalSales
2	FROM OrderDetail od
3	JOIN Product p ON od.ProductID = p.ProductID
4	GROUP BY p.CarModel
5	ORDER BY TotalSales DESC;

Data Output	Messages	Notifications																																				
<div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> </div> <table> <tr> <th></th><th>carmodel character varying (100)</th><th>totalsales numeric</th></tr> <tr><td>1</td><td>Galant</td><td>7040006.14</td></tr> <tr><td>2</td><td>Grand Prix</td><td>6793496.12</td></tr> <tr><td>3</td><td>Camaro</td><td>6786994.68</td></tr> <tr><td>4</td><td>SL-Class</td><td>6226512.28</td></tr> <tr><td>5</td><td>C70</td><td>6199948.43</td></tr> <tr><td>6</td><td>Suburban 1500</td><td>6086810.14</td></tr> <tr><td>7</td><td>Century</td><td>5875109.31</td></tr> <tr><td>8</td><td>M3</td><td>5331517.94</td></tr> <tr><td>9</td><td>E-Class</td><td>5088724.43</td></tr> <tr><td>10</td><td>Accord</td><td>5062198.76</td></tr> <tr><td>11</td><td>CL-Class</td><td>4980881.38</td></tr> </table>		carmodel character varying (100)	totalsales numeric	1	Galant	7040006.14	2	Grand Prix	6793496.12	3	Camaro	6786994.68	4	SL-Class	6226512.28	5	C70	6199948.43	6	Suburban 1500	6086810.14	7	Century	5875109.31	8	M3	5331517.94	9	E-Class	5088724.43	10	Accord	5062198.76	11	CL-Class	4980881.38		
	carmodel character varying (100)	totalsales numeric																																				
1	Galant	7040006.14																																				
2	Grand Prix	6793496.12																																				
3	Camaro	6786994.68																																				
4	SL-Class	6226512.28																																				
5	C70	6199948.43																																				
6	Suburban 1500	6086810.14																																				
7	Century	5875109.31																																				
8	M3	5331517.94																																				
9	E-Class	5088724.43																																				
10	Accord	5062198.76																																				
11	CL-Class	4980881.38																																				

Query 9: Subquery to find suppliers who have products costing more than the average price of all products

Query	Query History
1	SELECT s.SupplierName, p.CarModel, p.CarPrice
2	FROM Product p
3	JOIN Supplier s ON p.SupplierID = s.SupplierID
4	WHERE p.CarPrice > (SELECT AVG(CarPrice) FROM Product)
5	ORDER BY p.CarPrice DESC;

Data Output	Messages	Notifications
+		
supplename	carmodel	carprice
character varying (255)	character varying (100)	numeric
1	Kwideo	Corolla
2	Mycat	Canyon
3	Yata	900
4	Twitterlist	Mustang
5	Feedfire	Justy
6	Tazz	5 Series
7	Voomm	CL-Class
8	Livefish	Ram 2500
9	Alnyx	Diamante
10	Jamia	Parisienne
11	InnoZ	Solara

Query 10: Complex query with multiple joins and a subquery to list all products sold in the last month, sorted by quantity sold

Query

Query History

```
1 SELECT p.CarMaker, p.CarModel, SUM(od.Quantity) AS TotalQuantity
2 FROM OrderDetail od
3 JOIN Orders o ON od.OrderID = o.OrderID
4 JOIN Product p ON od.ProductID = p.ProductID
5 WHERE o.OrderDate >= (
6     SELECT DATE_TRUNC('month', MAX(OrderDate))
7     FROM Orders
8 ) AND o.OrderDate < (
9     SELECT (DATE_TRUNC('month', MAX(OrderDate)) + INTERVAL '1 month')
10    FROM Orders
11 )
12 GROUP BY p.CarMaker, p.CarModel
13 ORDER BY TotalQuantity DESC;
```

14

Data Output

Messages

Explain X

Notifications

</

XVI. QUERY EXECUTION ANALYSIS

Query: The provided SQL query calculates total sales by car model and involves a complex join between the Product and OrderDetail tables, with a subsequent aggregation and sorting operation. The query's execution plan indicates the use of hash joins and aggregates, which are performance-intensive due to full table scans and lack of indexes.

Query

Query History

1

SELECT p.CarModel, SUM(od.Sales) AS TotalSales

2

FROM OrderDetail od

3

JOIN Product p ON od.ProductID = p.ProductID

4

GROUP BY p.CarModel

5

ORDER BY TotalSales DESC;

Data Output

Messages

Explain X

Notifications

Query

Query History

1

explain analyze

2

SELECT p.CarModel, SUM(od.Sales) AS TotalSales

3

FROM OrderDetail od

4

JOIN Product p ON od.ProductID = p.ProductID

5

GROUP BY p.CarModel

6

ORDER BY TotalSales DESC;

Data Output

Messages

Explain X

Notifications

QUERY PLAN

test

1

Sort (cost=75.93..77.11 rows=473 width=39) (actual time=1.556..1.571 rows=473 loops=1)

2

Sort Key: (sum(od.sales)) DESC

3

Sort Method: quicksort Memory: 54kB

4

-> HashAggregate (cost=49.00..54.91 rows=473 width=39) (actual time=1.323..1.452 rows=473 loops=1)

5

Group Key: p.carmodel

6

-> Hash Join (cost=26.62..44.64 rows=872 width=15) (actual time=0.298..0.756 rows=872 loops=1)

7

Hash Cond: (od.productid = p.productid)

8

-> Seq Scan on orderdetail od (cost=0.00..15.72 rows=872 width=12) (actual time=0.029..0.157 rows=872 loop=1)

9

-> Hash (cost=15.72..15.72 rows=872 width=11) (actual time=0.261..0.262 rows=872 loops=1)

10

Buckets: 1024 Batches: 1 Memory Usage: 47kB

11

-> Seq Scan on product p (cost=0.00..15.72 rows=872 width=11) (actual time=0.011..0.126 rows=872 loop=1)

12

Planning Time: 0.223 ms

13

Execution Time: 1.642 ms



To optimize this query, indexing is recommended on ProductID in both joined tables for faster joins, CarModel in the Product table to expedite group operations, and Sales in the OrderDetail table to quicken aggregation computations. These indexes are expected to reduce full table scans, enhance data retrieval speed, and improve overall query performance by enabling more efficient join algorithms and faster grouping and sorting operations.

XVII. WEB APPLICATION

In the development of our web application using Streamlit, we aimed to provide a robust tool that enables users to interactively explore customer and supplier data associated with specific car makes and models, drawing information from a PostgreSQL database. This application facilitates the querying of detailed information about suppliers and customers by allowing users to select a car maker and model from dynamically populated dropdown menus. The backend integration is handled through psycopg2, ensuring direct interaction with our database, which is managed securely by fetching credentials from environment variables to protect sensitive information.

Customers Suppliers

Find Suppliers by Car

Choose Car Maker

Choose an option

Choose Car Model

Choose an option

Search Suppliers

Customers Suppliers

Find Customers by Car

Choose Car Maker

Choose an option

Choose Car Model

Choose an option

Search Customers

We employed Streamlit's caching mechanism to optimize the responsiveness of the application. By caching database connections and data fetching functions, we significantly improved the application's efficiency, reducing the need to establish new database connections or repeat database queries for repeated requests. This not only enhances the user experience by speeding up data retrieval but also reduces the load on the server, leading to a more scalable solution.

Customers Suppliers

Find Suppliers by Car

Choose Car Maker

Audi

Choose Car Model

Q7

Search Suppliers

	suppliername	suppliercontactdetails	city	state	country
0	Kamba	321625923	Inglewood	California	United States
1	Geba	439179977	Tulsa	Oklahoma	United States

Customers Suppliers

Find Customers by Car

Choose Car Maker

Audi

Choose Car Model

S6

Search Customers

	customername	gender	phonenummer	city	state	country
0	Hadleigh Baptiste	Male	4789269762	Macon	Georgia	United States
1	Kirsten Shooobridge	Female	6192786699	San Diego	California	United States

Moreover, we focused on user-friendly design elements within the Streamlit framework. The interface is divided into tabs, making it straightforward for users to switch between viewing suppliers and customers. This separation aids in clarity and usability. Each interaction—selecting a car maker and model, followed by the initiation of a search—triggers a backend process that fetches and displays relevant data in a well-organized table format. This setup ensures that users, regardless of their technical background, can easily navigate and utilize the application to gain the insights they need. This approach emphasizes our commitment to creating intuitive, efficient, and secure applications that cater to real-world business needs.

References

Dataset: [Supply chain management for Car \(kaggle.com\)](https://www.kaggle.com/datasets/ashishpatel26/supply-chain-management-for-car)

Preprocessing Notebook:

<https://colab.research.google.com/drive/1Hl2ocgDHfR-BgauiosTObbUIVeFYhX2f?usp=sharing>

<https://docs.streamlit.io/develop/tutorials/databases/postgresql>